



COMMUNITY DAY

PUNE 2022



COMMUNITY DAY

PUNE 2022



Scale your applications with DynamoDB

Sridevi Murugayen


Director – Fidelity Investments
AWS Data Hero
AWS User Group- Chennai

Disclaimer: The views or opinions expressed in this presentation are solely those of the speaker and do not necessarily represent those of Fidelity Investments. This research does not reflect in any way procedures, processes or policies of operations within Fidelity Investments.



COMMUNITY DAY

PUNE 2022





COMMUNITY DAY

PUNE 2022

AWS Database Services



**Amazon
RDS**

**Managed Relational
Database Service**
Postgres, MariaDB,
Oracle, MySQL & SQL
Server



**Amazon
DocumentDB**

**Fully MongoDB
Compatible
Document
Database**



**Amazon
ElastiCache**

**In-Memory Key Value
Store**



**Amazon
DynamoDB**

**Fully Managed
Key-Value and
Document Database**



**Amazon
Neptune**

**Fully Managed Graph
Database**



**Amazon
Aurora**

**Cloud-Native
Relational Database**



**Amazon
Timestream**

**Fully Managed
Time Series
Database**



**Amazon
QLDB**

**Fully Managed
Ledger Database**



**Amazon
MemoryDB
for Redis**

**Fully managed,
Redis-compatible,
in-memory
database service**



**Amazon
Keyspaces**

**Apache
Cassandra-co
mpatible
database
service**



COMMUNITY DAY

PUNE 2022



SQL vs NoSQL

SQL	NoSQL
Optimized for Storage	Optimized for Compute
Structured and Strict adherence, Table, Data types etc. (Normalized)	Distributed and non relational (no foreign keys)
Vertical scaling (Add memory, disk , CPU as required)	Horizontal Scaling
Designed for Random Access Pattern	Designed for Strict Access Pattern
Fetch criteria is not critical for DB design. Query is an after thought	Access pattern should be defined before we design the table



COMMUNITY DAY

PUNE 2022

Amazon DynamoDB

- Key-Value and Document Database
- Query Driven
- Support eventual & strong consistency
- Can handle One trillion requests per day with single-digit ms latency



Fully
managed



Consistently fast at any Scale



Secure



Integrates with AWS
Lambda,
Amazon Redshift, and
more



Cost-effective



COMMUNITY DAY

PUNE 2022

Journey of SnapChat

- 300 million daily active users
- 5 Billion+ snaps per day
- 400 TB of data
- 900+ EKS Clusters running

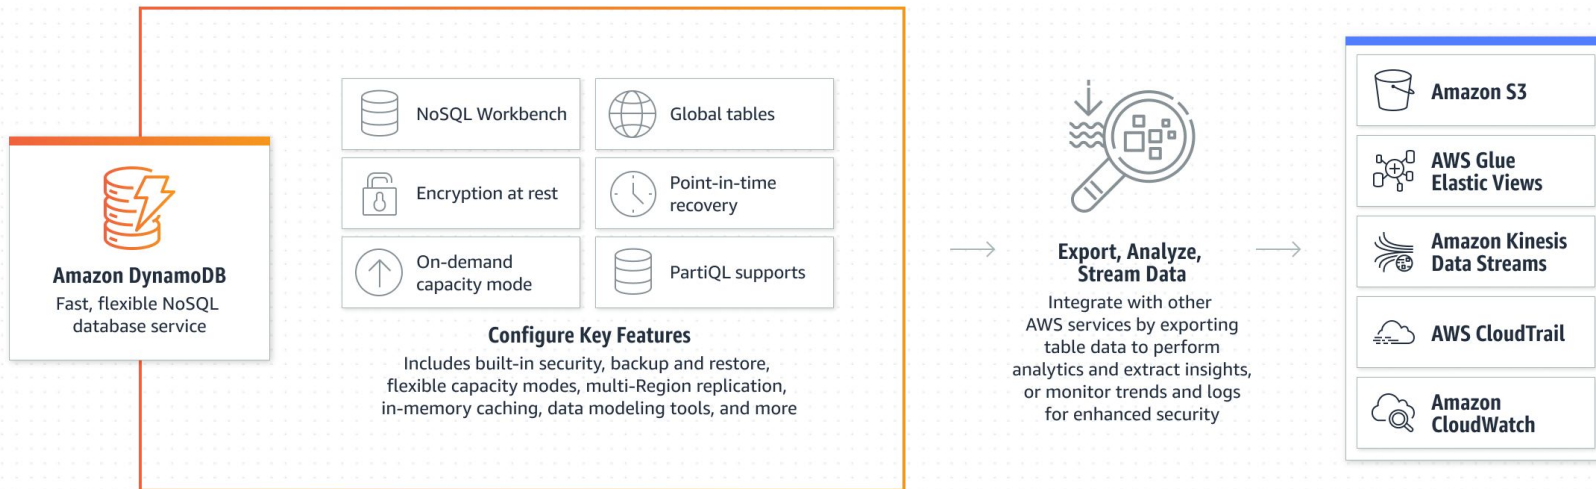
https://www.youtube.com/watch?v=Cgv0kfp_6xQ



COMMUNITY DAY

PUNE 2022

How DynamoDB works





COMMUNITY DAY

PUNE 2022



Understanding DynamoDB




COMMUNITY DAY

PUNE 2022

DynamoDB components

Table



Mandatory
Key-value access pattern
Determines data distribution

Tables

- DynamoDB stores data in tables.
- A table is a collection of data.

Items

- Each table contains zero or more items.
- An item is a group of attributes that is uniquely identifiable
- Maximum item size in DynamoDB is 400 KB

Attributes

- Each item is composed of one or more attributes.
- An attribute is a fundamental data element



COMMUNITY DAY

PUNE 2022



DynamoDB Primary Key

Simple Primary Key (Partition key)

- Composed of one attribute - partition key.
- Partition key value is input to an internal hash function.
- The output from the hash function determines the partition in which the item will be stored.
- No two items can have the same partition key value.

Composite Primary Key (Partition key and sort key)

- Composite primary key, composed of two attributes. the partition key, and the sort key.
- Partition key value is input to an internal hash function.
- The output from the hash function determines the partition in which the item will be stored.
- Possible for two items to have the same partition key value. However, those two items must have different sort key values.




COMMUNITY DAY

PUNE 2022

DynamoDB – Horizontal scale


- All data is stored in tables , must have a partition key
- A table of hypothetical customers

custid	name	occupation		
121375	Alejandro Rosalez	musician	instruments: [guitar, violin]	band: "Brave Bards"
743192	Mary Major	educator	area: physics	qual: PhD
615278	Nikki Wolf	pilot	level: ATP	type: [B777-200, B767]
415692	Zhang Wei	pilot	level: commercial	type: [PA28, C172, C152]



Storage node

custid	name
615278	Nikki Wolf



Storage node

custid	name
121375	Alejandro Rosalez
743192	Mary Major



Storage node

custid	name
415692	Zhang Wei





COMMUNITY DAY

PUNE 2022

Highly available and durable

OrderId: 1
CustomerId: 1
ASIN: [B00X4WHP5E]

hash(1) = 7B



Availability Zone A



Host 1



Host 2



Host 3

Availability Zone B



Host 4




Host 5



Host 6

Availability Zone C



Host 7



Host 8



Host 9

3-way replication


Data is always replicated to three Availability Zones



COMMUNITY DAY

PUNE 2022

DynamoDB – High-level architecture



Clients

Regional endpoint

Load balancers

Request routers

Storage nodes



© 2022, Amazon Web Services, Inc. or its affiliates. All rights reserved.



COMMUNITY DAY

PUNE 2022



Provisioning DynamoDB

WCU (Write Capacity Units)

- 1 write capacity unit (WCU) = 1 write of up to 1 KB/s.
- 2 WCUs = 1 transactional write request (one write per second) for items up to 1 KB.
- For writes greater than 1 KB, total number of writes required = $(\text{total item size} / 1 \text{ KB})$ rounded up

RCU (Read Capacity Units)

- 1 read capacity unit (RCU) = 1 strongly consistent read of up to 4 KB/s or 2 eventually consistent reads of up to 4 KB/s per read.
- 2 RCUs = 1 transactional read request (one read per second) for items up to 4 KB.
- For reads on items greater than 4 KB, total number of reads required = $(\text{total item size} / 4 \text{ KB})$ rounded up.



COMMUNITY DAY

PUNE 2022

DynamoDB Indices

Global secondary index

- An index with a partition key and sort key that can be different from those on the base table.
- Can be created after the base table is created
- Separate RCU and WCU to be provisioned

Local secondary index

- An index that has the same partition key as the table, but a different sort key.
- Should be created during table creation
- Uses the RCU and WCU of base table

Each table in DynamoDB has a quota of 20 global secondary indexes (default quota) and 5 local secondary indexes per table.

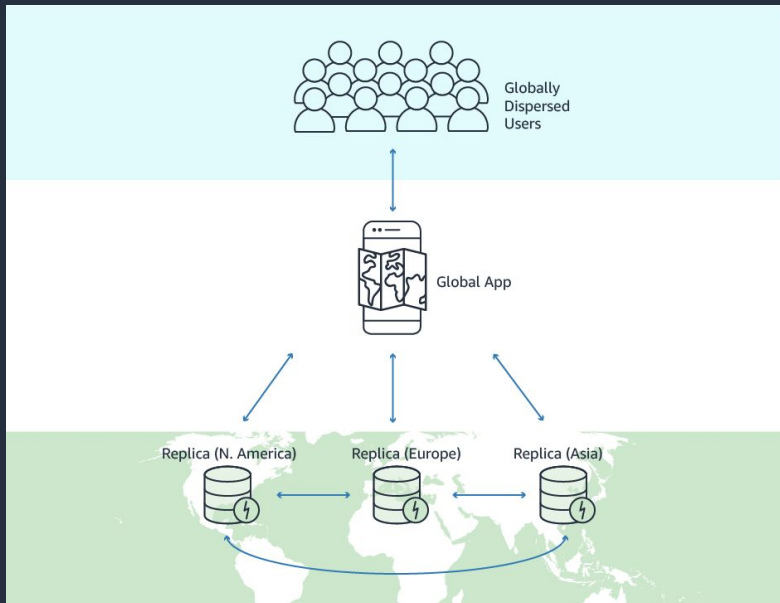


COMMUNITY DAY

PUNE 2022

Global Tables

fully-managed, multi-master, multi-region, multi active database



- Read and write locally, access your data globally
- single-digit-millisecond latency for your globally distributed application
- Replication happens for all attributes and no projection of attributes



COMMUNITY DAY

PUNE 2022

Data Modeling Process

- Create your Entity relationship diagram
- List out the access patterns (read and write access patterns)
- Design your primary keys to handle your access patterns
- Define Indices to facilitate additional access patterns



COMMUNITY DAY

PUNE 2022

Demo





COMMUNITY DAY

PUNE 2022



References

- <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/Introduction.html>
- <https://aws.amazon.com/dynamodb/global-tables/>
- https://explore.skillbuilder.aws/learn/course/internal/view/elearning/67/amazon-dynamodb-for-serverless-architectures?trk=980219ed-943e-411a-b498-769eb04b9884&sc_channel=el
- <https://www.youtube.com/watch?v=SC-YAPgJpms>
- <https://www.youtube.com/watch?v=yvBR71D0nAQ>
- <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/SampleData.CreateTables.html>



COMMUNITY DAY

PUNE 2022

Q&A





COMMUNITY DAY

PUNE 2022



Thank you!

See you at the AWS Community Day Pune 2023

