



COMMUNITY DAY

PUNE 2022



COMMUNITY DAY

PUNE 2022



Building application with analytical capabilities using AWS services and Apache Druid



COMMUNITY DAY

PUNE 2022

Before we start...

How many of you heard of Druid

How many of you are tried Druid in the past

How many of you are using / planning to use druid in production

How many of you have the use case of providing analytics tool to your users as part of the product offerings



COMMUNITY DAY

PUNE 2022

Again we are not starting ..

How many of you are application developers

How many of you are Data Scientist

How many of you are Data Engineers





COMMUNITY DAY

PUNE 2022

Who Am I



Tijo Thomas



Lead Solutions Architect @ Imply



Dealing with Big Data Challenges since 2014



[tijothomas](#)



@tijothomas214



tijothomas



COMMUNITY DAY

PUNE 2022



Agenda

- Story of building an analytical app
- Demo - Wiring it up with AWS Services
- Demo on performance and concurrency
- How is Druid is best fit analytical apps
- Q&A



COMMUNITY DAY

PUNE 2022

I want to build...

“ In my application landing page, I want to show a
real-time dashboard”



@tijothomas214



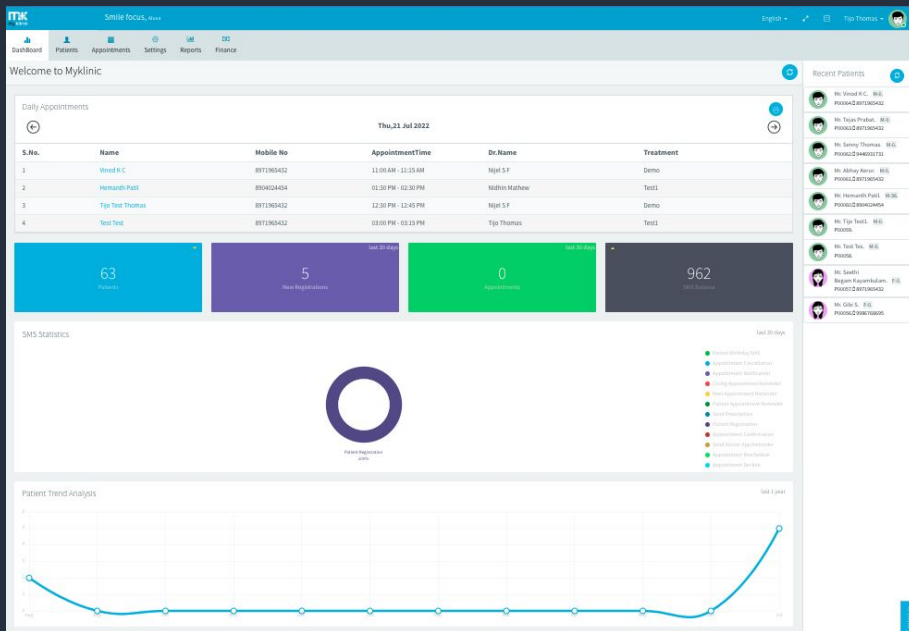
tijopthomas



COMMUNITY DAY

PUNE 2022

I want to build dashboard



Pain Areas

- This dashboard aggregate data over wide range of data points
- Dashboard refreshes frequent interval or when user navigate to landing page
- SaaS application with high concurrency
- Latency expectation is <200ms
- Realtime and historical data



@tjiothomas214



tjiothomas



COMMUNITY DAY

PUNE 2022

I want to build...

“ Wish to provide a self service reporting capability to my customers”



@tijothomas214



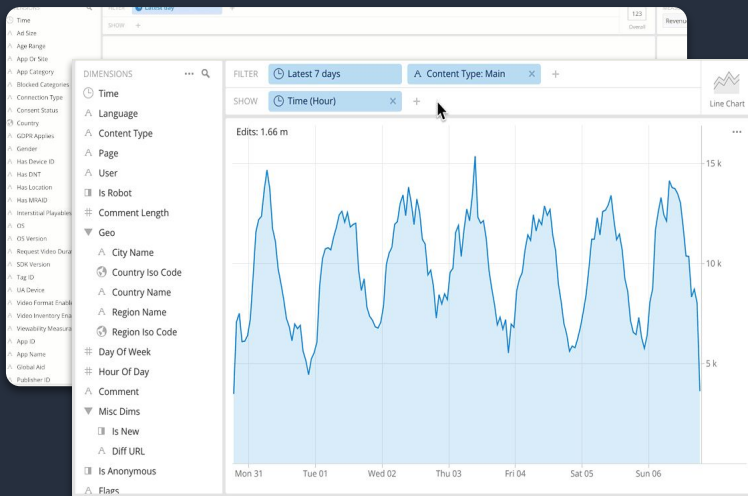
tijopthomas



COMMUNITY DAY

PUNE 2022

This is what a modern analytics app looks like



An interactive data experience with...

- TBs to PBs of data processed
- 100s to 1000s of queries per sec
- Real-time and historical data analyzed
- Most of the queries are groupby in nature
- Latency requirement < 200ms
- Highly concurrent query workload






COMMUNITY DAY

PUNE 2022

Sample Architecture ..



@tijothomas214



tijothomas



COMMUNITY DAY

PUNE 2022

Sample Architecture ..



@tijothomas214



tijopthomas



COMMUNITY DAY

PUNE 2022

Challenges

- Data warehouses including column stores work only with batch ingestion and are not optimal for streaming data.
- DWs are difficult to scale up and down based on the Query load.
- DWs support low concurrency. Concurrency comes with exponential cost
- DWs have limited support for indexing the data



@tijothomas214



tijophomas



COMMUNITY DAY

PUNE 2022



Demo onboarding data



@tijothomas214



tijopthomas



COMMUNITY DAY

PUNE 2022



Demo how to query Druid



@tijothomas214



tijopthomas



COMMUNITY DAY

PUNE 2022



Demo the performance and concurrency



@tijothomas214




tijopthomas



COMMUNITY DAY

PUNE 2022

Unlimited Users: Modern analytics apps drive up concurrency



**The cost for high
concurrency increases
exponentially with CDWs**

A different database is needed for high concurrency at the best value



COMMUNITY DAY

PUNE 2022

Druid shines for modern analytics applications



	Best Fit		
	 General Purpose DBs	 Cloud Data Warehouses	 Apache Druid
Interactive Analytics at Any Scale	-	-	Yes
High Concurrency (at the Best Value)	-	Yes, but expensive	Yes
Insights on Streaming and Batch Data	-	-	Yes



COMMUNITY DAY

PUNE 2022

Recap the requirements

I want to build ...

- High concurrency analytical application
- Wanted to analyze real time and historical data together
- Need latency < 200ms for most of the usecase
- Wanted to scale up or down the infrastructure within reasonable time frame
eg: <1 hr
- Many other secondary requirements



@tijothomas214



tijophomas



COMMUNITY DAY

PUNE 2022

Secondary requirement

Reduce cost per query over older data with higher latency and lesser concurrency

Dont wanted to ingest all raw data but wanted to pre aggregate while ingesting to DB



@tijothomas214



tijopthomas



COMMUNITY DAY

PUNE 2022



Let us see how druid helps here



@tijothomas214



tijopthomas

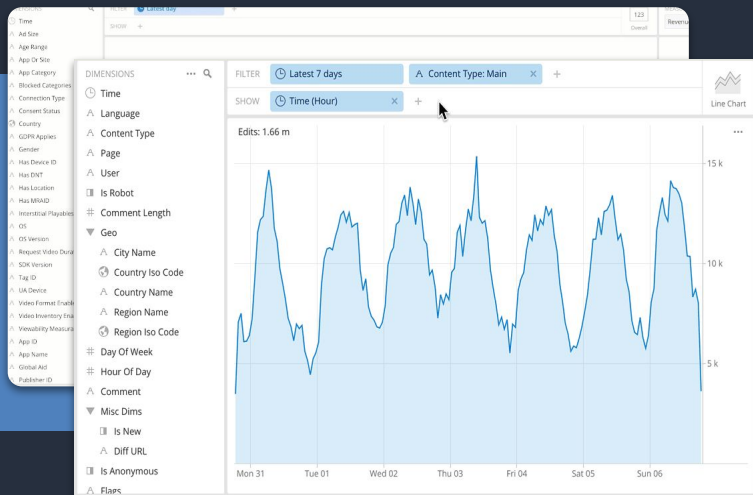


COMMUNITY DAY

PUNE 2022



What Is Druid?



Druid is the best database
for **modern analytics**
applications

10+
Years

400+
Contributors

1,000+
Customers deployed

11,000+
Community members



COMMUNITY DAY

PUNE 2022



druid[®] is...

High performance.

low-latency, distributed query execution and high throughput ingestion

Real-time.

event data (clickstream, network flows, user behavior, programmatic advertising, server metrics, IoT...)

Analytics.

counting, ranking, statistics...

Database.

highly-available, time-sharded, partitioned, columnar, indexed, compressed, versioned materialized view



COMMUNITY DAY

PUNE 2022



Commonly asked questions

Druid a replacement for our existing Data warehouse ?

Where does Druid fit in my big data stack?



@tijothomas214



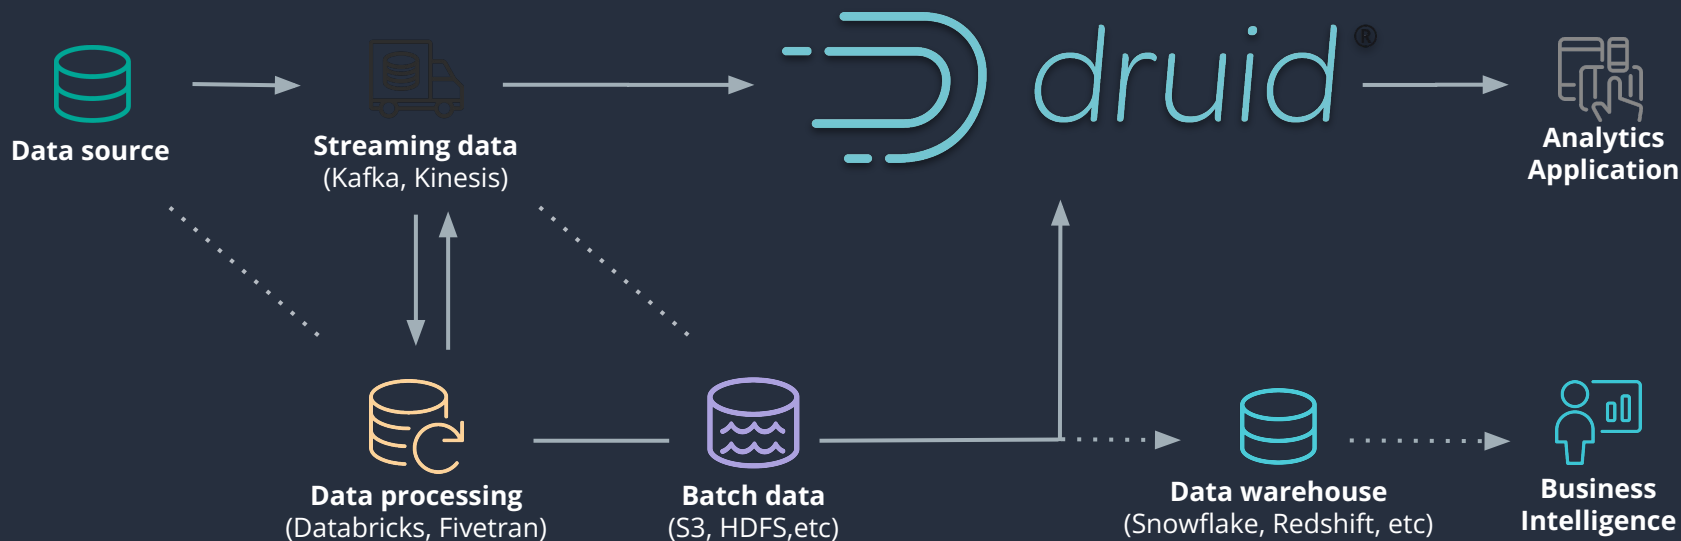
tijopthomas



COMMUNITY DAY

PUNE 2022

Druid in the Data Ecosystem





COMMUNITY DAY

PUNE 2022

Commonly asked questions

Is Druid a log aggregation search system



@tijothomas214



tijopthomas



COMMUNITY DAY

PUNE 2022

Commonly asked questions

Is Druid a in-memory database?



@tijothomas214



tijopthomas




COMMUNITY DAY

PUNE 2022

Is Druid a in-memory database?

Query Processing

- Segments can be memory mapped
- Not every time druid need to read from disk






COMMUNITY DAY

PUNE 2022

Is Druid a in-memory database?



- While ingesting data will be kept in memory and then written to disk
- The threshold for this is mentioned in the ingestion specification




@tijothomas214



tijophomas



imply &



druid[®]



Imply Community Team

community@imply.io



Imply Training Program

<https://learn.imply.io>



Druid User Forum by Imply

<https://www.druidforum.org>



Druid Professionals Group

<https://www.linkedin.com/groups/8791983/>



Druid Meetups

<https://www.meetup.com/pro/apache-druid/>



Druid News & Info

@druidio #apachedruid @implydata



Druid Slack

#druid

<https://druid.apache.org/community/join-slack>



Druid @ Implies

<https://imply.io/what-is-druid/>



COMMUNITY DAY

PUNE 2022

Q&A





COMMUNITY DAY

PUNE 2022



Thank you!

See you at the AWS Community Day Pune 2023

